

Linear & Ridge Regression Analysis

Data and Libraries

```
In [3]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Ridge
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import GridSearchCV
```

```
In [4]: # Import historical almond pricing and production data
almond_df = pd.read_excel('AlmondData.xlsx')

# Import historical US population data
uspopulation_df = pd.read_excel('USPopulation.xlsx')

# Import historical USD Index data
usindex_df = pd.read_excel('USDollarIndex.xlsx')
```

Data Preparation

```
In [6]: # Merge US population data into almond data frame
almond_df['USPopulation_Millions'] = uspopulation_df['July 1 (million people)1']
```

```
In [7]: # Filter US Index data from 1981 to 2024
usindex_1981to2024 = usindex_df.loc[usindex_df['Year (last business day)'] >= 1981, :]

# Reset index
usindex_1981to2024 = usindex_1981to2024.reset_index()
```

```
In [8]: # Merge USDIndex data into almond data frame
almond_df['USDIndex'] = usindex_1981to2024['DXY close']
```

```
In [9]: # Save marketing years
marketing_year = almond_df['Marketing_Year']
```

```
In [10]: # Drop unnecessary columns
uncolumns = ['Marketing_Year', 'Yield_Lbs_Per_Acre',
             'Value_of_Production_Usd_Mill', 'Utilized_Production_Lbs_Mill',
             'Current_Reserve_Lbs_Mill', 'Per capita_Availability_Lbs']

# Drop columns
almond_df.drop(columns=uncolumns, inplace=True)
```

```
In [11]: # Show final data frame
almond_df.head()
```

```
Out[11]:
```

	Bearing_Acres_Th	Total_Production_Lbs_Mill	Price_Cents_Per_Lb	Loss_Production_Lbs_Mill	Marketable_P
0	326.8	322	147.0	16.86	
1	326.2	408	78.0	24.87	
2	339.0	347	94.0	16.24	
3	360.0	242	104.0	20.21	
4	381.0	590	77.4	26.36	

Linear Regression Analysis

```
In [13]: # Create X feature matrix
X = almond_df.drop(columns='Price_Cents_Per_Lb')

# Select target variable
y = almond_df['Price_Cents_Per_Lb']
```

```
In [14]: # Split data into training and testing sets (80%/20%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

```
In [15]: # Create model
model = LinearRegression()

# Fit model
model.fit(X_train, y_train)
```

```
Out[15]:
```

▼ LinearRegression

LinearRegression()

```
In [16]: # Test model
y_pred = model.predict(X_test)
```

```
In [17]: # Calculate error metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R² Score: {r2:.4f}")
```

Mean Absolute Error (MAE): 32.71
Mean Squared Error (MSE): 2473.63
Root Mean Squared Error (RMSE): 49.74
R² Score: 0.7233

```
In [18]: # Get the feature importance (coefficients)
coefficients = pd.DataFrame({'Feature': X.columns, 'Coefficient': model.coef_})

# Print coefficients
print(coefficients.sort_values(by='Coefficient', ascending=False))
```

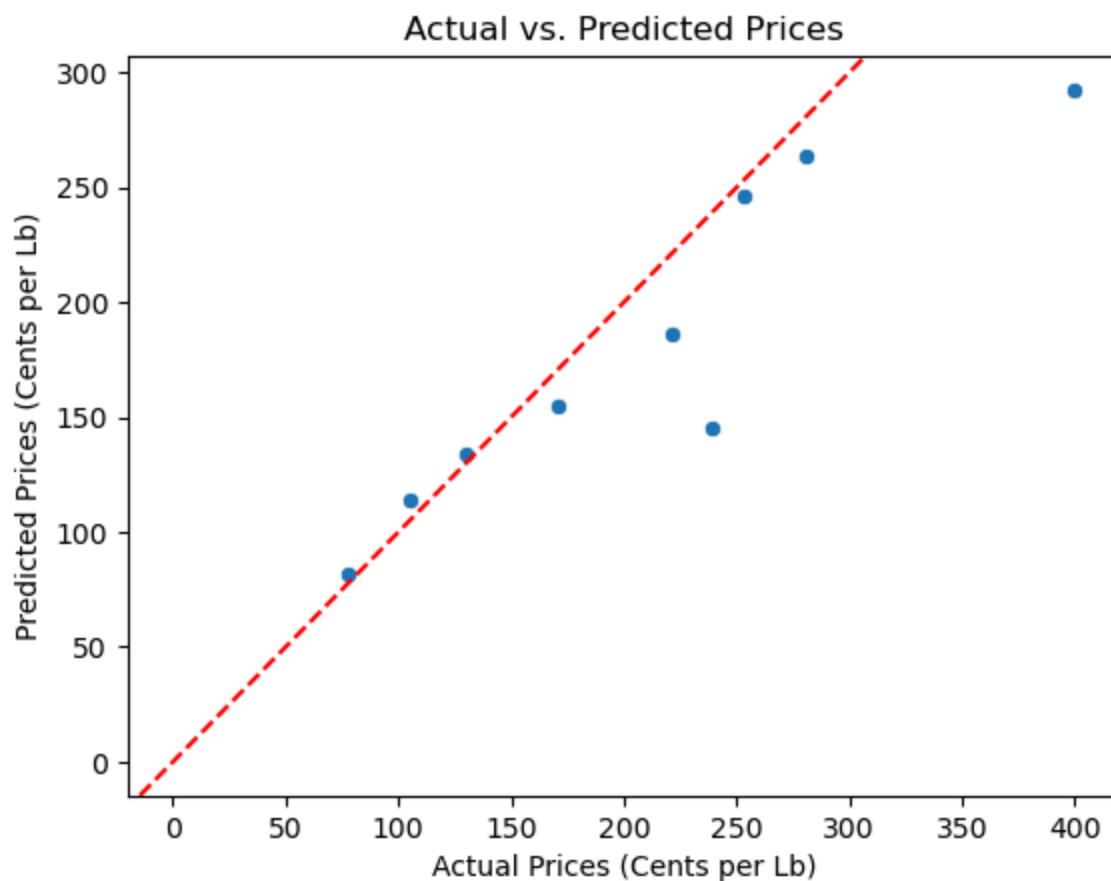
	Feature	Coefficient
4	Imports_lbs_Mill	3.967334
6	Supply_Lbs_Mill	2.122925
8	Exports_Lbs_Mill	0.846494
7	Ending_Stocks_Lbs_Mill	0.774407
10	USPopulation_Millions	0.728366
2	Loss_Production_Lbs_Mill	0.675250
9	Domestic_Availability_Lbs_Mill	0.502024
0	Bearing_Acres_Th	0.279754
5	Beginning_Stocks_Lbs_Mill	-0.003078
11	USDIndex	-1.023544
1	Total_Production_Lbs_Mill	-1.169156
3	Marketable_Production_Lbs_Mill	-1.844406

```
In [19]: # Plot scatter plot
sns.scatterplot(x=y_test, y=y_pred)

# Show ideal line
plt.axline([0, 0], [1, 1], color="red", linestyle="--")

# Plot Labels
plt.xlabel("Actual Prices (Cents per Lb)")
plt.ylabel("Predicted Prices (Cents per Lb)")
plt.title("Actual vs. Predicted Prices")

# Show chart
plt.show()
```



Ridge Regression Analysis

```
In [21]: # Create empty data frame
vif_data = pd.DataFrame()

# Add feature matrix as columns
vif_data["Feature"] = X.columns

# Compute the Variance Inflation Factor (VIF) for each feature
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]

# Show VIF
print(vif_data)
```

	Feature	VIF
0	Bearing_Acres_Th	353.929946
1	Total_Production_Lbs_Mill	inf
2	Loss_Production_Lbs_Mill	inf
3	Marketable_Production_Lbs_Mill	inf
4	Imports_lbs_Mill	23.959837
5	Beginning_Stocks_Lbs_Mill	6262.688105
6	Supply_Lbs_Mill	inf
7	Ending_Stocks_Lbs_Mill	inf
8	Exports_Lbs_Mill	inf
9	Domestic_Availability_Lbs_Mill	inf
10	USPopulation_Millions	78.593529
11	USDIndex	50.163597

```
C:\Users\emili\anaconda3\Lib\site-packages\statsmodels\stats\outliers_influence.py:198: RuntimeWarning: divide by zero encountered in scalar divide
  vif = 1. / (1. - r_squared_i)
```

```
In [22]: # Create standard scaler
scaler = StandardScaler()

# Fit scaler
X_train_scaled = scaler.fit_transform(X_train)

# Transform test matrix
X_test_scaled = scaler.transform(X_test)
```

```
In [23]: # Initialize Ridge regression
ridge_model = Ridge(alpha=1.0)

# Fit training data into model
ridge_model.fit(X_train_scaled, y_train)

# Make predictions
y_pred = ridge_model.predict(X_test_scaled)
```

```
In [24]: # Calculate error metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.4f}")
```

Mean Absolute Error (MAE): 31.10
Mean Squared Error (MSE): 2595.78
Root Mean Squared Error (RMSE): 50.95
R² Score: 0.7097

```
In [25]: # Get the feature importance (coefficients)
coefficients = pd.DataFrame({'Feature': X.columns, 'Coefficient': ridge_model.coef_})

# Show coefficients
print(coefficients.sort_values(by='Coefficient', ascending=False))
```

	Feature	Coefficient
4	Imports_lbs_Mill	63.297927
0	Bearing_Acres_Th	27.864404
10	USPopulation_Millions	20.523656
8	Exports_Lbs_Mill	-0.277304
2	Loss_Production_Lbs_Mill	-3.913615
5	Beginning_Stocks_Lbs_Mill	-6.271804
6	Supply_Lbs_Mill	-8.053045
1	Total_Production_Lbs_Mill	-9.090824
3	Marketable_Production_Lbs_Mill	-9.170580
7	Ending_Stocks_Lbs_Mill	-11.888543
11	USDIndex	-14.186478
9	Domestic_Availability_Lbs_Mill	-22.994026

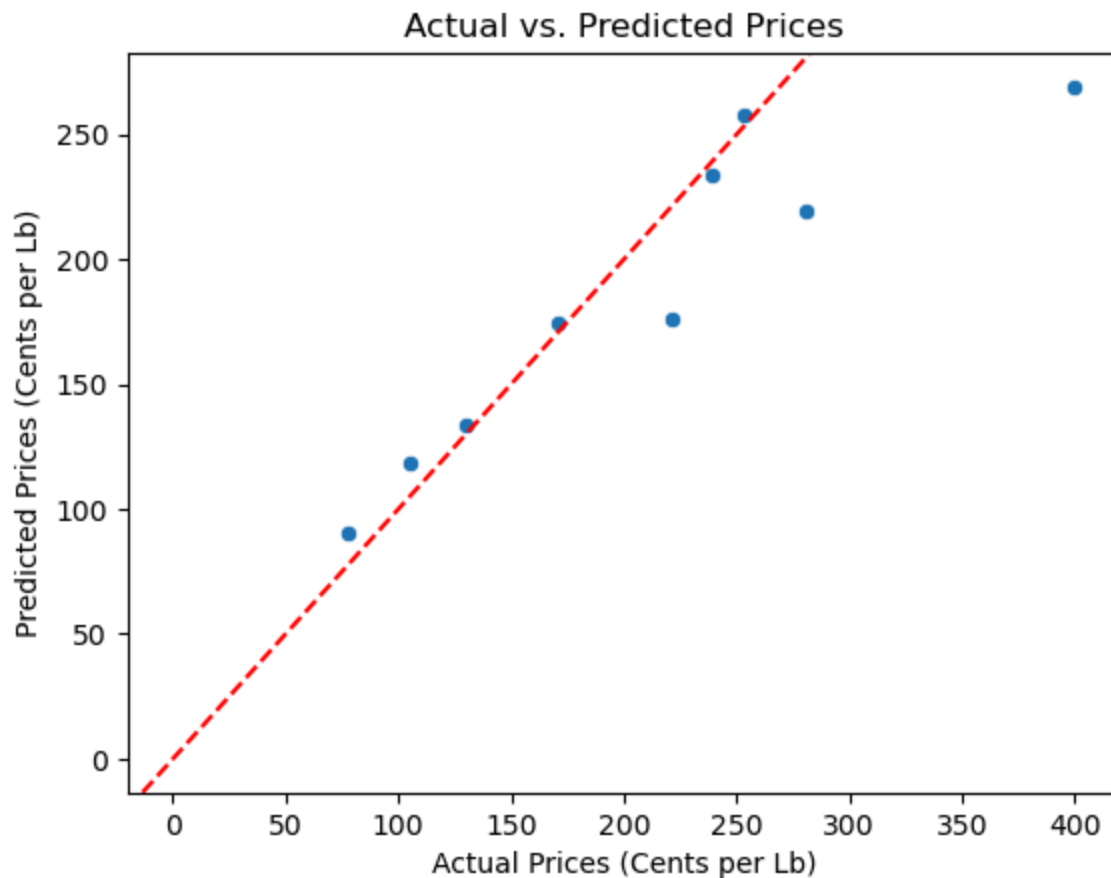
```
In [26]: # Plot scatter plot
sns.scatterplot(x=y_test, y=y_pred)

# Show ideal line
```

```
plt.axline([0, 0], [1, 1], color="red", linestyle="--")

# Plot labels
plt.xlabel("Actual Prices (Cents per Lb)")
plt.ylabel("Predicted Prices (Cents per Lb)")
plt.title("Actual vs. Predicted Prices")

# Show chart
plt.show()
```



Ridge Model Optimization

```
In [28]: # Define a range of alpha values
alphas = {'alpha': np.logspace(-3, 3, 10)}
```

```
In [29]: # Create Grid Search model with Cross-Validation
ridge_cv = GridSearchCV(Ridge(), alphas, cv=5, scoring='r2')

# Fit Grid Search model
ridge_cv.fit(X_train_scaled, y_train)
```

```
Out[29]:
```

▸ **GridSearchCV**

▸ **estimator: Ridge**

▸ Ridge

```
In [30]: # Get the best alpha
best_alpha = ridge_cv.best_params_['alpha']
```

```
# Print result
print(f"Best Alpha: {best_alpha}")
```

Best Alpha: 2.154434690031882

```
In [31]: # Set Ridge model with optimized alpha value
ridge_model_optimal = Ridge(alpha=best_alpha)

# Fit training data into Ridge model
ridge_model_optimal.fit(X_train_scaled, y_train)
```

```
Out[31]: ▼ Ridge
Ridge(alpha=2.154434690031882)
```

```
In [32]: # Make predictions with optimized Ridge model
y_pred_optimal = ridge_model_optimal.predict(X_test_scaled)
```

```
In [33]: # Calculate error metrics
mae = mean_absolute_error(y_test, y_pred_optimal)
mse = mean_squared_error(y_test, y_pred_optimal)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred_optimal)

# Print evaluation metrics
print(f"Optimized Mean Absolute Error (MAE): {mae:.2f}")
print(f"Optimized Mean Squared Error (MSE): {mse:.2f}")
print(f"Optimized Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"Optimized R² Score: {r2:.4f}")
```

Optimized Mean Absolute Error (MAE): 36.26
Optimized Mean Squared Error (MSE): 3233.40
Optimized Root Mean Squared Error (RMSE): 56.86
Optimized R² Score: 0.6383

```
In [34]: # Get the feature importance (coefficients)
coefficients = pd.DataFrame({'Feature': X.columns, 'Coefficient': ridge_model_optimal.coef_})

# Show coefficients
print(coefficients.sort_values(by='Coefficient', ascending=False))
```

	Feature	Coefficient
4	Imports_lbs_Mill	54.560576
10	USPopulation_Millions	16.154383
0	Bearing_Acres_Th	15.279725
8	Exports_Lbs_Mill	-0.141693
2	Loss_Production_Lbs_Mill	-3.505363
5	Beginning_Stocks_Lbs_Mill	-4.517457
6	Supply_Lbs_Mill	-4.977698
1	Total_Production_Lbs_Mill	-5.722350
3	Marketable_Production_Lbs_Mill	-5.751424
9	Domestic_Availability_Lbs_Mill	-9.885883
7	Ending_Stocks_Lbs_Mill	-12.762904
11	USDIndex	-13.444154

```
In [35]: # Plot optimized Ridge regression model
sns.scatterplot(x=y_test, y=y_pred_optimal)
```

```
# Show ideal line
plt.axline([0, 0], [1, 1], color="red", linestyle="--")

# Plot labels
plt.xlabel("Actual Prices (Cents per Lb)")
plt.ylabel("Predicted Prices (Cents per Lb)")
plt.title("Optimized Ridge Model - Actual vs. Predicted Prices")

# Show chart
plt.show()
```

